

Abdulmalik Alquwayfli

☎ 055 476 7376 ✉ af.alquwayfli@gmail.com 📍 Riyadh 💼 Abdulmalik Alquwayfli 🌐 AbdulmalikDS
🔗 Portofolio

PROFILE

Driven to create real impact, I specialize in advanced deeptech solutions leveraging Natural Language Processing (NLP), Large Language Models (LLMs), and Agentic AI. I apply research-backed methods to design intelligent systems that simplify complex problems, enhance decision-making, and deliver measurable business value.

EDUCATION

Reinforcement Learning, CS234, Stanford University Nov 2025

Bachelor of Information Systems, Imam Mohammed Bin Saud Universty Sep 2025

EXPERIENCE

Artificial Intelligence Developer, TAHAKOM • NLP Department 📁 Jul 2025 – Present

- Architected and deployed continuous **pre-training and post-training** pipelines on **AWS SageMaker Hyperpod**, leveraging Trainium, **GPU clusters**, and Hugging Face Optimum Neuron, covering **SFT, DPO, RLHF** for advanced alignment.
- Engineered a **parameter-efficient fine-tuning** pipeline for the **Qwen-2.5-72B** model using **LoRA** and **DeepSpeed ZeRO-3** on AWS A100, creating a **bilingual** instruction-following assistant specialized in **Saudi dialects**.
- Developed **LangGraph**-based **autonomous agents** with dynamic **tool calling** and real **code execution**, integrating **deep search** and file management. Collaborated with cross-functional teams including **Data Science, Product, and Video Understanding** on multiple use cases.
- Designed and **orchestrated 5+ autonomous agents** via **MCPs**, integrating ArXiv and GitHub APIs to **automate literature** review by **searching papers, code and data**.
- Conducted deep **literature reviews** on **agentic AI frameworks** (SICA, **LangGraph**, ReAct, Camel AI, Pydantic AI) and analyzed technical reports of state-of-the-art **LLM architectures** and **alignment strategies**.

Machine learning Enginner, Muragib 📁 Mar 2025 – May 2025

- **Trained, fine-tuned**, and evaluated 10+ open-source **BERT**-based models (e.g., CAMEL BERT, AraBERT), **achieving over 90% accuracy** with <60ms latency, optimized for real-time applications.
- Developed and deployed advanced machine learning models for a wide range of **NLP tasks**, including **sentiment analysis**, sarcasm detection, location distribution analysis.
- Built scalable, **AI-powered analytics** systems that deliver actionable insights from unstructured text, empowering business stakeholders to make fast, data- driven decisions.

Research And Development Engineer, واعي | Wae 📁 Nov 2024 – Mar 2025

- Designed, developed, and deployed AI-based solutions for hands-on workshops and internal tools, **directly impacting over 150 employees and 50+ students**.
- Contributed to the **design and delivery** of educational and technical content, bridging the gap between AI theory and practice in professional settings
- **Researched the latest AI tools** and technologies to drive innovation.

Software Developer - Artificial Intelligence VW, Microsoft Jun 2024 – Jul 2024

- Built an AI-powered **computer vision model** using Custom Vision to **recognize and classify images**.
- Conducted **data preprocessing and feature engineering** for predictive analytics solutions.

SKILLS

Programming languages: Python, Java, JavaScript, SQL, Cuda, PHP, Rust, C++, R

Libraries, tools, and frameworks: Hugging Face, LLaMA-Factory, Unsloth, vLLM, LangGraph, LangChain, PyTorch, Keras, Scikit-learn, NLTK, OpenCV, Pandas, NumPy, Matplotlib, Seaborn, FastAPI, Ollama, Pydantic, wandb, Temporal

DevOps & Environments: Git, Bash, Linux, Ubuntu, Docker, Conda, Jupyter Notebook, Google Collaboration, Postman, Curls, Cursor, AWS SageMaker AI, HyperPod, HPC, SLURM, PyTorch Lightning, LogFire

Databases: PostgreSQL, MongoDB, Qdrant, ChromaDB, Metabase, Supabase, DBeaver, Alembic, Ssqlalchemy

Cloud: Digital Ocean, AWS (SageMaker, Trainium), Alibaba Cloud, TPU, AI Foundry, Microsoft Azure

Model Architectures: Attention, Transformers, BERT, RoBERTa, DDP, SFT, DPO ,GRPO, RLHF, GPT, RAG, LoRA, Quantization, Agentic AI, Multi-Agent Systems, MCPS, MOE, MOA, RNNs EfficientNet, ConvNeXt

COURSES & CERTIFICATIONS

AWS AI & ML Scholars , <i>AWS & Udacity</i> ↗	Jun 2025
Building AI Browser Agents , <i>AGI, inc</i> ↗	Apr 2025
LLMOps: Building Real-World Applications With LLMs , <i>Udacity</i> ↗	Dec 2024
Finetuning Large Language Models , <i>Lamini</i> ↗	Dec 2024
Open Source Models , <i>Hugging Face</i> ↗	Dec 2024
Neural Networks and Deep Learning , <i>DeepLearning.AI</i> ↗	Sep 2024
Full Stack Deep Learning , <i>The Ful Stack</i>	Dec 2024
Python for Data Science, AI & Development , <i>IBM</i> ↗	Jun 2024

PROJECTS

Quantigence: Multi-Agent AI Framework for Quantum Security	May 2025 – Jun 2025
Independent research on post-quantum cryptographic threats using a multi-agent system coordinating five quantized LLMS (Qwen2.5, Mistral, Llama3.1, Phi-3, Gemma2). Runs on consumer GPUs with Q4_K_M quantization and CUDA acceleration. Integrates data from Brave Search, arXiv, GitHub, Puppeteer, and SQLite using Multi-Context-Protocol (MCPS)	
Farqad The Pioneer in Smart Financial Conversations , <i>GP</i> ↗	Dec 2024 – May 2025
Farqad addresses the challenge of simplifying financial management and advice for Arabic and English-speaking users by developing a bilingual RAG chatbot with state of the art models, highly engineered and maintainable code	
PerceiveAI , <i>AI League Finalist</i> ↗	Feb 2025 – Apr 2025
This project aimed at making sports commentary accessible to the deaf and hard- of-hearing. By leveraging cutting-edge AI technologies, PerceiveAI translates live sports commentary into text then to sign language using animated avatars	
langgraph-runtime-inmem-open , <i>LangGraph Contribution</i> ↗	Jul 2025
open-source LangGraph runtime tools for agent memory management	
Multi-Agent System Visualizer ↗	Jul 2025 – Sep 2025
A sophisticated 3D visualization platform for multi-agent research systems with real-time AI-powered research capabilities.	
Arabian Sheep Classification , <i>Kaggle Computer Vision Competition</i> ↗	Jun 2025
Scored First in the leaderboard Day One and developed AI system for classifying 7 Arabian sheep breeds celebrating Eid al-Adha heritage, achieving championship-level 96.67% macro F1 score using ConvNeXt deep learning architecture with advanced data augmentation and Test-Time Augmentation techniques.	

PUBLICATIONS

Quantigence: A Multi-Agent AI Framework for Quantum Security Research ↗	Jun 2025
--	----------